

# Airbnb listings in Paris, France\*

Xu Qi

March 4, 2024

## Table of contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Data and Processing</b>	<b>1</b>
<b>3</b>	<b>Results</b>	<b>2</b>
3.1	Distribution and properties of individual variables . . . . .	2
3.2	Relationships between variables . . . . .	9
	<b>References</b>	<b>11</b>

## 1 Introduction

This paper is a exploratory data analysis of Airbnb listings in Paris, France, as at 12 December 2023.

## 2 Data and Processing

The dataset is from Inside Airbnb (Cox 2021). Data was collected and analyzed using the statistical programming software R (R Core Team 2023), with additional support packages including `tidyverse` (Wickham et al. 2019), `ggplot2` (Wickham 2016), `dplyr` (Wickham et al. 2023), `janitor` (Firke 2023), `knitr` (Xie 2014) , `naniar` (Tierney and Cook 2023) , `patchwork` (Pedersen 2024) , `arrow` (Richardson et al. 2024), `here` (Müller and Bryan 2020) and `modelsummary` (Arel-Bundock 2022).

---

\*Code and data from this analysis are available at:<https://github.com/xuqi2002/Paris-airbnb>

### 3 Results

#### 3.1 Distribution and properties of individual variables

Figure 1.1 Distribution of prices




Figure 1.2 Using the log scale for prices more than \$1,000




Figure 1: Distribution of prices of Paris Airbnb rentals in December 2023

We can look at the distribution of prices (Figure 1.1). There are outliers, so again we might like to consider it on the log scale (Figure 1.2).

Figure 2.1 Prices less than \$1,000 suggest some bunching




Figure 2.2 Prices between \$90 and \$210 illustrate the bunching more clearly




Figure 2: Distribution of prices for Airbnb listings in Paris in December 2023

Turning to Figure 2, if we focus on prices that are less than 1,000, then we see that most properties have a nightly price less than 250 (Figure 2 .1). It might be that this is happening around numbers ending in zero or nine. Let us just zoom in on prices between \$90 and \$210, out of interest, but change the bins to be smaller (Figure 2 .2).




Figure 3: Distribution of review scores rating for Paris Airbnb rentals in December 2023

On Airbnb, guests can give one to five star ratings across a variety of different aspects, including cleanliness, accuracy, value, and others. But when we look at the reviews in our dataset, it is clear that it is effectively a binary, and almost entirely the case that either the rating is five stars or not (Figure 3).




Figure 4: Distribution of review scores for Paris Airbnb rentals in December 2023

14519 properties do not have a review rating yet because they do not have enough reviews. It is a large proportion of the total so we might like to look at this in more detail using counts. We are interested to see whether there is something systematic happening with these properties. For instance, if the NAs were being driven by, say, some requirement of a minimum number of reviews, then we would expect they would all be missing.

One approach would be to just focus on those that are not missing and the main review score (Figure 4).




Figure 5: Distribution of review scores for properties with NA response time, for Paris Airbnb rentals in December 2023

There is an issue with NAs as there are a lot of them. For instance, we might be interested to see if there is a relationship with the review score (Figure 5). There are a lot that have an overall review of 100.




Figure 6: Missing values in Paris Airbnb data, by host response time

Usually missing values are dropped by `ggplot2` (Wickham 2016). We can use `geom_miss_point()` from `naniar` (Tierney and Cook 2023) to include them in the graph (Figure 6).




Figure 7: Distribution of the number of properties a host has on Airbnb, for Paris Airbnb rentals in December 2023

We might be interested in how many properties a host has on Airbnb (Figure 7). Based on Figure Figure 7 we can see there are a large number who have somewhere in the 100 to 500 properties range, with the usual long tail. The number with that many listings is unexpected and worth following up on. And there are a bunch with NA that we will need to deal with.



### 3.2 Relationships between variables




Figure 8: Relationship between price and review and whether a host is a superhost, for Paris Airbnb rentals in December 2023

We can look at the relationship between price and reviews, and whether they are a super-host, for properties with more than one review (Figure 8).

Table 1

	(1)
(Intercept)	−16.262 (0.481)
host_response_timewithin a day	2.019 (0.211)
host_response_timewithin a few hours	2.695 (0.210)
host_response_timewithin an hour	2.972 (0.209)
review_scores_rating	2.624 (0.089)
Num.Obs.	22 047
AIC	24 165.0
BIC	24 205.0
Log.Lik.	−12 077.507
RMSE	0.43

From Table 1, We see that each of the levels is positively associated with the probability of being a superhost. However, having a host that responds within an hour is associated with individuals that are superhosts in our dataset

## References

- Arel-Bundock, Vincent. 2022. “modelssummary: Data and Model Summaries in R.” *Journal of Statistical Software* 103 (1): 1–23. <https://doi.org/10.18637/jss.v103.i01>.
- Cox, Murray. 2021. “Get the Data.” *Insideairbnb.com*. <http://insideairbnb.com/get-the-data.html>.
- Firke, Sam. 2023. *Janitor: Simple Tools for Examining and Cleaning Dirty Data*. <https://CRAN.R-project.org/package=janitor>.
- Müller, Kirill, and Jennifer Bryan. 2020. *Here: A Simpler Way to Find Your Files*. <https://CRAN.R-project.org/package=here>.
- Pedersen, Thomas Lin. 2024. *Patchwork: The Composer of Plots*. <https://CRAN.R-project.org/package=patchwork>.
- R Core Team. 2023. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Richardson, Neal, Ian Cook, Nic Crane, Dewey Dunnington, Romain François, Jonathan Keane, Dragoş Moldovan-Grünfeld, Jeroen Ooms, Jacob Wujciak-Jens, and Apache Arrow. 2024. *Arrow: Integration to 'Apache' 'Arrow'*. <https://github.com/apache/arrow/>.
- Tierney, Nicholas, and Dianne Cook. 2023. “Expanding Tidy Data Principles to Facilitate Missing Data Exploration, Visualization and Assessment of Imputations.” *Journal of Statistical Software* 105 (7): 1–31. <https://doi.org/10.18637/jss.v105.i07>.
- Wickham, Hadley. 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D’Agostino McGowan, Romain François, Garrett Grolemund, et al. 2019. “Welcome to the tidyverse.” *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Wickham, Hadley, Romain François, Lionel Henry, and Kirill Müller. 2023. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.
- Xie, Yihui. 2014. *Knitr: A Comprehensive Tool for Reproducible Research in R*. Edited by Victoria Stodden, Friedrich Leisch, and Roger D. Peng. Chapman; Hall/CRC. <http://www.crcpress.com/product/isbn/9781466561595>.